

Boston Housing: KNN; Bias-Variance Trade-Off; Cross Validation

Chicago Booth ML Team

OVERVIEW

This R Markdown script uses the *Boston Housing* data set to illustrate the following:

- The **k-Nearest Neighbors (KNN)** algorithm;
- The **Bias-Variance Trade-Off**; and
- The use of **Cross Validation** to estimate Out-of-Sample (OOS) prediction error and determine optimal hyper-parameters, in this case the number of nearest neighbors k .

first, some boring logistics...

Let's first load some necessary R packages and helper functions and set the random number generator's seed:

```
# load CRAN libraries from CRAN packages
library(data.table)
library(ggplot2)
library(kknn)
# load modules from the common HelpR repo
helpr_repo_raw_url <- 'https://raw.githubusercontent.com/ChicagoBoothML/HelpR/master'
source(file.path(helpr_repo_raw_url, 'docv.R')) # this has docvknn used below

# set randomizer's seed
set.seed(99) # Gretzky was #99
```

Boston Housing Data Set

Let's then look at the **Boston Housing** data set:

```
# download data and read data into data.table format
boston_housing <- fread(
  'https://raw.githubusercontent.com/ChicagoBoothML/DATA__BostonHousing/master/BostonHousing.csv')
# count number of samples
nb_samples <- nrow(boston_housing)
# sort data set by increasing lstat
setkey(boston_housing, lstat)
boston_housing
```

##		crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio
##	1:	1.46336	0	19.58	0	0.605	7.489	90.8	1.9709	5	403	14.7
##	2:	1.83377	0	19.58	1	0.605	7.802	98.2	2.0407	5	403	14.7
##	3:	0.03359	75	2.95	0	0.428	7.024	15.8	5.4011	3	252	18.3
##	4:	0.57529	0	6.20	0	0.507	8.337	73.3	3.8384	8	307	17.4
##	5:	0.08664	45	3.44	0	0.437	7.178	26.3	6.4798	5	398	15.2
##	---											
##	502:	18.81100	0	18.10	0	0.597	4.628	100.0	1.5539	24	666	20.2
##	503:	1.62864	0	21.89	0	0.624	5.019	100.0	1.4394	4	437	21.2
##	504:	11.10810	0	18.10	0	0.668	4.906	100.0	1.1742	24	666	20.2
##	505:	45.74610	0	18.10	0	0.693	4.519	100.0	1.6582	24	666	20.2
##	506:	18.49820	0	18.10	0	0.668	4.138	100.0	1.1370	24	666	20.2

```
##      black lstat medv
## 1: 374.43  1.73 50.0
## 2: 389.61  1.92 50.0
## 3: 395.62  1.98 34.9
## 4: 385.91  2.47 41.7
## 5: 390.49  2.87 36.4
## ---
## 502:  28.79 34.37 17.9
## 503: 396.90 34.41 14.4
## 504: 396.90 34.77 13.8
## 505:  88.27 36.98  7.0
## 506: 396.90 37.97 13.8
```

This data set has **506** samples.

We'll focus on **using the *lstat* variable to predict the *medv* variable**. Let's first plot them against each other:

```
plot_boston_housing_data <- function(boston_housing_data,
                                     title='Boston Housing: medv vs. lstat',
                                     plot_predicted=TRUE) {
  g <- ggplot(boston_housing_data) +
    geom_point(aes(x=lstat, y=medv, color='actual'), size=2) +
    ggtitle(title) +
    xlab('medv') + ylab('lstat')


  if (plot_predicted) {
    g <- g +
      geom_line(aes(x=lstat, y=predicted_medv, color='predicted'), size=0.6) +
      scale_colour_manual(name='medv',
                          values=c(actual='blue', predicted='darkorange'))
  } else {
    g <- g +
      scale_colour_manual(name='medv',
                          values=c(actual='blue'))
  }

  g <- g +
    theme(plot.title=element_text(face='bold', size=24),
          axis.title=element_text(face='italic', size=18))

  g
}

plot_boston_housing_data(boston_housing, plot_predicted=FALSE)
```

Boston Housing: medv vs. lstat



k -Nearest Neighbors algorithm and Bias-Variance Trade-Off

```
try_k <- 5
```

Let's now try fitting a KNN predictor, with $k = 5$, of *medv* from *lstat*, using the entire 506 samples:


```
knn_model <- knn(medv ~ lstat,
  train=boston_housing, test=boston_housing[, .(lstat)],
  k=try_k, kernel='rectangular')
boston_housing[, predicted_medv := knn_model$fitted.values]
```

```
##      crim  zn  indus  chas   nox   rm   age   dis rad tax ptratio
##  1:  1.46336  0 19.58    0 0.605 7.489 90.8 1.9709  5 403    14.7
##  2:  1.83377  0 19.58    1 0.605 7.802 98.2 2.0407  5 403    14.7
##  3:  0.03359 75  2.95    0 0.428 7.024 15.8 5.4011  3 252    18.3
##  4:  0.57529  0  6.20    0 0.507 8.337 73.3 3.8384  8 307    17.4
##  5:  0.08664 45  3.44    0 0.437 7.178 26.3 6.4798  5 398    15.2
## ---
## 502: 18.81100  0 18.10    0 0.597 4.628 100.0 1.5539 24 666    20.2
## 503:  1.62864  0 21.89    0 0.624 5.019 100.0 1.4394  4 437    21.2
## 504: 11.10810  0 18.10    0 0.668 4.906 100.0 1.1742 24 666    20.2
## 505: 45.74610  0 18.10    0 0.693 4.519 100.0 1.6582 24 666    20.2
## 506: 18.49820  0 18.10    0 0.668 4.138 100.0 1.1370 24 666    20.2
##      black lstat medv predicted_medv
##  1: 374.43  1.73 50.0          42.60
##  2: 389.61  1.92 50.0          42.60
##  3: 395.62  1.98 34.9          42.60
##  4: 385.91  2.47 41.7          42.30
##  5: 390.49  2.87 36.4          43.96
## ---
## 502:  28.79 34.37 17.9          12.38
## 503: 396.90 34.41 14.4          12.38
```

```
## 504: 396.90 34.77 13.8          12.30
## 505:  88.27 36.98  7.0          13.38
## 506: 396.90 37.97 13.8          13.38
```

```
plot_boston_housing_data(boston_housing, title=paste('KNN Model with k =', try_k))
```

KNN Model with k = 5



With $k = 5$ – a small number of nearest neighbors – we have a very “squiggly” predictor, which fits the training data well but is **over-sensitive to small changes** in the *lstat* variable. We call this a **LOW-BIAS, HIGH-VARIANCE** predictor. We don’t like it.

```
try_k <- 200
```

Now, with, say, $k = 200$, we have the following:


```
knn_model <- kkn(boston_housing[, 1:11],
  train=boston_housing[, 1:11], test=boston_housing[, 12],
  k=try_k, kernel='rectangular')
boston_housing[, predicted_medv := knn_model$fitted.values]
```

```
##      crim  zn  indus  chas   nox    rm   age   dis rad tax ptratio
##  1:  1.46336  0 19.58    0 0.605 7.489 90.8 1.9709  5 403    14.7
##  2:  1.83377  0 19.58    1 0.605 7.802 98.2 2.0407  5 403    14.7
##  3:  0.03359 75  2.95    0 0.428 7.024 15.8 5.4011  3 252    18.3
##  4:  0.57529  0  6.20    0 0.507 8.337 73.3 3.8384  8 307    17.4
##  5:  0.08664 45  3.44    0 0.437 7.178 26.3 6.4798  5 398    15.2
##  ---
## 502: 18.81100  0 18.10    0 0.597 4.628 100.0 1.5539 24 666    20.2
```

```
## 503: 1.62864 0 21.89 0 0.624 5.019 100.0 1.4394 4 437 21.2
## 504: 11.10810 0 18.10 0 0.668 4.906 100.0 1.1742 24 666 20.2
## 505: 45.74610 0 18.10 0 0.693 4.519 100.0 1.6582 24 666 20.2
## 506: 18.49820 0 18.10 0 0.668 4.138 100.0 1.1370 24 666 20.2
##      black lstat medv predicted_medv
## 1: 374.43 1.73 50.0 29.9000
## 2: 389.61 1.92 50.0 29.9000
## 3: 395.62 1.98 34.9 29.9000
## 4: 385.91 2.47 41.7 29.9000
## 5: 390.49 2.87 36.4 29.9000
## ---
## 502: 28.79 34.37 17.9 15.6425
## 503: 396.90 34.41 14.4 15.6425
## 504: 396.90 34.77 13.8 15.6425
## 505: 88.27 36.98 7.0 15.6425
## 506: 396.90 37.97 13.8 15.6425
```

```
plot_boston_housing_data(boston_housing, title=paste('KNN Model with k =', try_k))
```

KNN Model with k = 200



Meh..., we're not exactly jumping around with joy with this one, either. The predictor line is **not over-sensitive**, but too "smooth" and too simple, **not responding sufficiently to significant changes in *lstat***. We call this a **HIGH-BIAS, LOW-VARIANCE** predictor.

```
try_k <- 50
```


Let's try something in between, say, $k = 50$, to see if we have any better luck:

```
knn_model <- kknnc(medv ~ lstat,
  train=boston_housing, test=boston_housing[, .(lstat)],
  k=try_k, kernel='rectangular')
boston_housing[, predicted_medv := knn_model$fitted.values]
```

```
##      crim zn indus chas  nox   rm   age   dis rad tax ptratio
##  1:  1.46336 0 19.58    0 0.605 7.489 90.8 1.9709  5 403    14.7
##  2:  1.83377 0 19.58    1 0.605 7.802 98.2 2.0407  5 403    14.7
##  3:  0.03359 75  2.95    0 0.428 7.024 15.8 5.4011  3 252    18.3
##  4:  0.57529 0  6.20    0 0.507 8.337 73.3 3.8384  8 307    17.4
##  5:  0.08664 45  3.44    0 0.437 7.178 26.3 6.4798  5 398    15.2
##  ---
## 502: 18.81100 0 18.10    0 0.597 4.628 100.0 1.5539 24 666    20.2
## 503:  1.62864 0 21.89    0 0.624 5.019 100.0 1.4394  4 437    21.2
## 504: 11.10810 0 18.10    0 0.668 4.906 100.0 1.1742 24 666    20.2
## 505: 45.74610 0 18.10    0 0.693 4.519 100.0 1.6582 24 666    20.2
## 506: 18.49820 0 18.10    0 0.668 4.138 100.0 1.1370 24 666    20.2
##      black lstat medv predicted_medv
##  1: 374.43  1.73 50.0          39.718
##  2: 389.61  1.92 50.0          39.718
##  3: 395.62  1.98 34.9          39.718
##  4: 385.91  2.47 41.7          39.718
##  5: 390.49  2.87 36.4          39.718
##  ---
## 502:  28.79 34.37 17.9          11.934
## 503: 396.90 34.41 14.4          11.934
## 504: 396.90 34.77 13.8          11.934
## 505:  88.27 36.98  7.0          11.934
## 506: 396.90 37.97 13.8          11.934
```

```
plot_boston_housing_data(boston_housing, title=paste('KNN Model with k =', try_k))
```

KNN Model with $k = 50$



Now, this looks pretty reasonable, and we'd think this predictor would **generalize well** when facing new, not yet seen, data. This is a **low-bias, low-variance** predictor. We love ones like this.

Hence, the key take-away is that, throughout a range of **hyper-parameter** k from small to large, we have seen a spectrum of corresponding predictors from “low-bias high-variance” to “high-bias low-variance”. This phenomenon is called the **BIAS-VARIANCE TRADE OFF**, a fundamental concept in Machine Learning that is applicable to not only KNN alone but to all modeling methods.

The bias-variance trade-off concerns the **generalizability of a trained predictor** in light of new data it's not seen before. If a predictor has high bias and/or high variance, it will not do well in new cases. **Good, generalizable predictors** need to have **both low bias and low variance**.

Out-of-Sample Error and Cross-Validation

To **quantify the generalizability of a predictor**, we need to estimate its **out-of-sample (OOS) error**, i.e. a certain measure of **how well the predictor performs on data not used in its training process**.

A popular way to produce such OOS error estimates is to perform **cross validation**. Refer to lecture slides or here for discussions on cross validation.

```
NB_CROSS_VALIDATION_FOLDS <- 5
NB_CROSS_VALIDATIONS <- 6
```

Now, let's consider **Root Mean Square Error (RMSE)** as our predictor-goodness evaluation criterion and use 5-fold cross validation 6 times to pick a KNN predictor that has satisfactory RMSE.

```

k_range = 2 : 200
cross_validations_rmse = data.table(k=k_range, cv_avg_rmse=0.)
for (i in 1 : NB_CROSS_VALIDATIONS) {
  this_cross_validation_rmse =
    sqrt(docvknn(boston_housing[, .(lstat)], boston_housing$medv,
                 k=k_range, nfold=NB_CROSS_VALIDATION_FOLDS,
                 verbose=FALSE) / nb_samples)
  cross_validations_rmse[, (paste('cv_', i, '_rmse', sep=''))] =
    this_cross_validation_rmse
  cross_validations_rmse[, cv_avg_rmse := cv_avg_rmse +
    (this_cross_validation_rmse - cv_avg_rmse) / i]
}

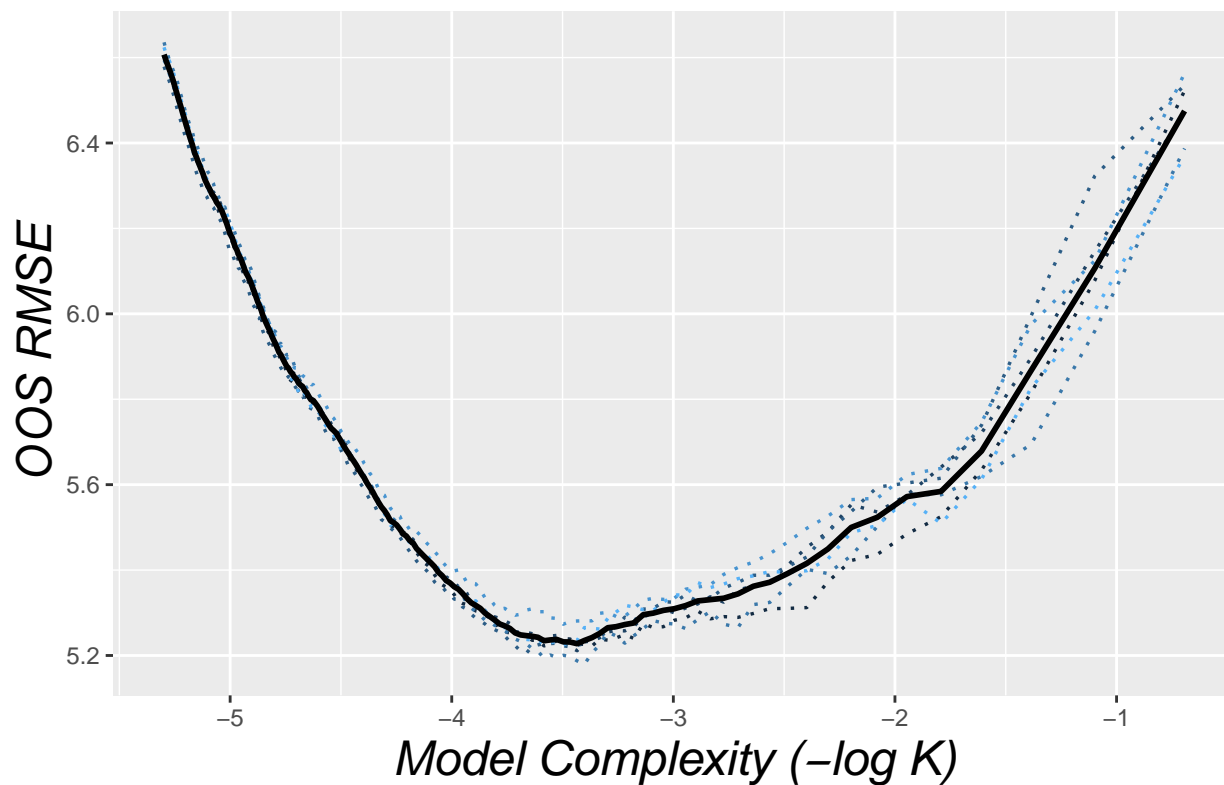
g <- ggplot(cross_validations_rmse)

for (i in 1 : NB_CROSS_VALIDATIONS) {
  g <- g + geom_line(aes_string(x='-log(k)', y=(paste('cv_', i, '_rmse', sep='')),
                                color=i, linetype='dotted', size=0.6))
}

g <- g +
  geom_line(aes(x=-log(k), y=cv_avg_rmse),
            color='black', size=1) +
  ggtitle('Cross Validations') +
  xlab('Model Complexity (-log K)') + ylab('OOS RMSE') +
  guides(color=FALSE) +
  theme(plot.title=element_text(face='bold', size=24),
        axis.title=element_text(face='italic', size=18))
g

```

Cross Validations




```
best_k = k_range[which.min(cross_validations_rmse$cv_avg_rmse)]
```


From the above plot, the best k , one that minimizes the average cross-validation RMSE, is **31**, which produces the following predictor:

```
knn_model <- kknnc(medv ~ lstat,
                   train=boston_housing, test=boston_housing[, .(lstat)],
                   k=best_k, kernel='rectangular')
boston_housing[, predicted_medv := knn_model$fitted.values]
```

```
##      crim zn indus chas  nox   rm   age   dis rad tax ptratio
##  1:  1.46336 0 19.58    0 0.605 7.489 90.8 1.9709  5 403    14.7
##  2:  1.83377 0 19.58    1 0.605 7.802 98.2 2.0407  5 403    14.7
##  3:  0.03359 75  2.95    0 0.428 7.024 15.8 5.4011  3 252    18.3
##  4:  0.57529 0  6.20    0 0.507 8.337 73.3 3.8384  8 307    17.4
##  5:  0.08664 45  3.44    0 0.437 7.178 26.3 6.4798  5 398    15.2
##  ---
## 502: 18.81100 0 18.10    0 0.597 4.628 100.0 1.5539 24 666    20.2
## 503:  1.62864 0 21.89    0 0.624 5.019 100.0 1.4394  4 437    21.2
## 504: 11.10810 0 18.10    0 0.668 4.906 100.0 1.1742 24 666    20.2
## 505: 45.74610 0 18.10    0 0.693 4.519 100.0 1.6582 24 666    20.2
## 506: 18.49820 0 18.10    0 0.668 4.138 100.0 1.1370 24 666    20.2
##      black lstat medv predicted_medv
##  1: 374.43  1.73 50.0      42.03548
##  2: 389.61  1.92 50.0      42.03548
##  3: 395.62  1.98 34.9      42.03548
##  4: 385.91  2.47 41.7      42.03548
##  5: 390.49  2.87 36.4      41.44194
##  ---
## 502:  28.79 34.37 17.9      11.83548
## 503: 396.90 34.41 14.4      11.83548
## 504: 396.90 34.77 13.8      11.83548
## 505:  88.27 36.98  7.0      11.83548
## 506: 396.90 37.97 13.8      11.83548
```

```
plot_boston_housing_data(boston_housing, title=paste('KNN Model with k =', best_k))
```

KNN Model with k = 31



BONUS: implementation by the *caret* package

caret is a popular R package that provides standardized interfaces with 200+ Machine Learning algorithms.

Much of the above procedures can be re-done very succinctly with *caret* as follows:

```
library(caret)

cross_validated_knn_model =
  train(medv ~ lstat, data=boston_housing,
        method='kkn',
        tuneGrid=expand.grid(kmax=200,
                              kernel='rectangular',
                              distance=2),
        trControl=trainControl(method='repeatedcv',
                                number=NB_CROSS_VALIDATION_FOLDS,
                                repeats=NB_CROSS_VALIDATIONS,
                                allowParallel=TRUE))

cross_validated_knn_model

## k-Nearest Neighbors
##
## 506 samples
## 1 predictor
```

```
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 6 times)
## Summary of sample sizes: 405, 405, 405, 404, 405, 405, ...
## Resampling results:
##
##   RMSE      Rsquared
##  5.234369  0.6766338
##
## Tuning parameter 'kmax' was held constant at a value of 200
##
## Tuning parameter 'distance' was held constant at a value of 2
##
## Tuning parameter 'kernel' was held constant at a value of rectangular
##
best_k = cross_validated_knn_model$finalModel$best.parameters$k
```

The best k identified by *caret* is **39**. Note that there can be a range of acceptable “best” hyper-parameters because of randomization.